

AI-Native Sales Copilot — Onboarding PRD

Owner	Rohit Saraf, Product
Status	Draft · v1.0
Date	2026-06-11
Reviewers	Engineering · Design · GTM Leadership
Personas	RevOps Admin (Priya) · AE/AM/CSM Seller (Alex, Riley)

Purpose. This document specifies the end-to-end onboarding experience for the RGI Workbench — HG Insights' AI-native GTM Intelligence platform. It covers the **RevOps admin onboarding** (Steps 1–9: tenant context, offerings, scoring, workbench, plays, integrations, teams, agents, invitations) and the **seller first-run experience** (Step 10 + ongoing workspace). The design principle: the platform knows everything it needs from the tenant's domain — admins review and refine, not build from scratch.

■ Target: RevOps admin reaches a fully-scored workbench in <60 minutes. Seller is productive within 10 minutes of first login.

1. Strategic Context

Most GTM tools require RevOps to manually configure ICP filters, build segments, pick competitors, and set up scoring rules. Customers stall mid-setup and never reach value. The RGI Workbench flips this: **HG already knows everything needed to bootstrap a tenant from their domain alone.**

What HG knows	What it enables
Firmographic data on the tenant company	Pre-fill company profile, size, HQ, industry
Product + market intelligence on what the tenant sells	Auto-propose offerings with pain points + ICP
Install + intent data on competitors	Auto-generate scoring models + displacement signals
IT spend and AI spend per company	Score whitespace universe before admin enters the workbench

By the time the admin clicks *Enter Workbench*, the platform has already extracted tenant context, configured up to 3 offerings, generated a scoring model per offering (Fit / Need / Intent), scored the HG-universe account pool, and tagged everything with provenance. The onboarding is agentic from screen one.

2. Goals & Non-Goals

Goal	Detail
Agentic 3-step core (RevOps)	Admin completes tenant context → offerings → scoring in ~30 min
AI extracts, admin refines	Every field pre-populated from HG data + LLM. No blank-slate authoring.
Traceable + editable	Every AI-extracted value shows its source. Admin can override any field.
Max 3 offerings at onboarding	LLM groups products if tenant has more; extras added post-onboarding
Score before entry	Scoring model + initial HG-universe run completes before admin enters workbench
Session-persistent	Admin can close browser and resume exactly where they left off

Out of Scope (v1)

- Authoring custom scoring models from scratch (admin tunes the AI-generated model post-onboarding)
- Configuring Plays during onboarding (post-onboarding Workbench activity)
- More than 3 offerings at onboarding (hard cap; admin can add post-onboarding)
- Multi-language support (English-only, v1)
- Multi-tenant admin (single admin per tenant; can invite teammates post-onboarding)
- Integrations beyond Salesforce, HubSpot, Slack in v1

3. Personas

Persona	Role in Onboarding	Time Investment
RevOps Admin (Priya)	Drives all admin steps (1–9). Final say on context, offerings, scoring, plays + teaming agency invitations	30–60 min (agency onboarding)
AE / AM Seller (Alex, Riley)	Doesn't participate in admin onboarding. Enters at Step 10 via magic link	< 2 min setup + 30-sec tour
MOPs / Data Lead (Marcus)	Optional reviewer of scoring model in Step 3.	0–15 min if engaged

4. End-to-End Onboarding Flow

The onboarding is a sequential 10-step flow. Steps 1–9 are owned by the RevOps admin; Step 10 is triggered per seller when they accept their invitation. Every step is resumable across sessions. Completed steps can be revisited at any time from the Workbench.

STEP 1	RevOps	Tenant Context Extraction	Enter domain → AI extracts firmographics, products, ICP, competitors. Admin reviews + confirms.
STEP 2	RevOps	Offering Configuration	AI groups products into ≤3 offerings with pain points, intent topics, competitors, ICP. Admin confirms each.
STEP 3	RevOps	Scoring Models + Scoring	AI generates Fit/Need/Intent model per offering. System scores HG-universe accounts. Admin reviews tier distribution + confirms.
STEP 4	RevOps	Workbench Entry	Admin enters scored workbench: 7 default columns (Tier, Why Now, Competitive Insights...). Enrich with AI adds columns via natural language.
STEP 5	RevOps	Play Configuration	AI proposes 3–5 Plays (Competitive Takeout, Net New, Expansion, etc.) composed from 1P + 3P signals. Admin activates ≥1.
STEP 6	RevOps	Integrations Setup	Connect Salesforce, HubSpot, Slack. Field mapping, sync direction, agent write-back scopes (all off by default). Integration completion unlocks Step 7.
STEP 7	RevOps	Configure Teams	Create seller teams, each with offerings, a scoring profile, default Plays, and audience. Teams are the inheritance layer for seller config.
STEP 8	RevOps	Configure Agents	Enable seller-facing agents (Account Brief, Email Outreach, Opportunity Finder, Find More Contacts) at tenant/team/seller level.
STEP 9	RevOps	Invite Sellers	CRM-driven seller discovery. Auto-assign to teams. Bulk-invite via email + Slack DM. Sellers inherit team config.
STEP 10	Seller	Seller First-Run	Seller accepts magic link → account setup → workbench pre-loaded with book + whitespace + scores + Plays + agents. 5-step orientation tour.

5. Step 1 — Tenant Context Extraction

User Story: I enter my company's domain. The platform tells me what it already knows about us so I don't have to type it in. I review, fix what's wrong, and approve.

Extracted Context (AI-populated, all fields editable)

Group	Fields
Firmographics	Industry (NAICS), employee count, revenue, HQ, public status, year founded
Product / Market	Products[], pain points[], IT spend categories[], intent topics[]
Market Positioning	Competitors[], complementary tech[], buying committee personas[]
Target ICP	Industries (tier 1/2/3), geography, revenue band, employee band

Key Requirements

ID	Requirement
F1.1	Accept domain input; trigger Tenant Context Extraction agent.
F1.2	Agent uses HG firmographics + web search + LLM to populate all fields within 60s.
F1.3	Every field shows its source (e.g., 'HG firmographic database', 'Web · tenant.com/products').
F1.4	All fields are editable inline; edits captured in audit log (who, when).
F1.5	Low-confidence fields (< 0.6) are visually flagged 'Needs review'. Fields < 0.3 block confirmation.
F1.6	Admin can re-run extraction per field or for the entire context.
F1.7	Admin must explicitly confirm context before proceeding. Confirmation creates a versioned snapshot.
F1.8	If domain has no HG data, surface a manual-entry fallback (same fields, no pre-fill).
F1.9	Extraction is idempotent; re-run shows diff vs. prior extraction.

Acceptance Criteria

- Given a valid domain, all fields in §5 are extracted within 60s.
- Every extracted field has a visible source chip.
- Admin can edit any field; edits persist immediately and appear in audit log.
- Admin cannot proceed to Step 2 until 'Confirm tenant context' is clicked.
- Fields with confidence < 0.6 are visually flagged; < 0.3 blocks confirmation.
- If domain has no HG data, manual-entry fallback works for all fields.
- Re-extraction is idempotent and shows a diff vs. previous extraction.

6. Step 2 — Offering Configuration

User Story: The platform organizes my products into up to 3 offerings with the right competitors, ICP, and pain points. I review the grouping, adjust if needed, and confirm.

Offering Grouping Logic

If tenant has ≤ 3 products $\rightarrow$ each becomes its own offering. If $> 3 \rightarrow$ LLM groups into max 3 offerings by shared buyer persona, target ICP, and competitor overlap. The LLM surfaces a *grouping rationale* per offering so the admin can verify the logic.

Field	Source	Editable?
Name	LLM-suggested or product name	Yes
Description	LLM-generated	Yes
Products[]	Subset of tenant.products[]	Yes (drag to reassign)
Pain Points[]	LLM-filtered from tenant.painPoints[]	Yes (add/remove/edit)
Intent Topics[]	LLM-filtered from tenant.intentTopics[]	Yes
Competitors[]	LLM-filtered from tenant.competitors[]	Yes
Complementary Tech[]	LLM-filtered from tenant.complementaryTech[]	Yes
Target ICP	Inherits from tenant or LLM-narrowed per offering	Yes

Key Requirements

ID	Requirement
F2.1	Auto-generate 1–3 offerings on entry to Step 2, each with grouping rationale.
F2.2	Admin can drag a product between offerings; dependent attributes (pain points, competitors) re-evaluate.
F2.3	Hard cap of 3 offerings during onboarding (enforced in UI).
F2.4	Removing an offering re-homes its products to a holding bucket; admin must reassign.
F2.5	Admin must confirm each offering individually before proceeding to Step 3.
F2.6	All offering edits persist immediately.

Acceptance Criteria

- 1–3 offerings appear on Step 2 entry, each with grouping rationale.
- Admin can drag products between offerings; attributes re-evaluate.
- All offering attributes are editable and persist.
- Hard cap of 3 offerings is enforced.
- Admin cannot proceed to Step 3 until all offerings are confirmed.

7. Step 3 — Scoring Models + Initial Universe Scoring

User Story: After confirming offerings, the platform shows me a scoring model per offering and the top accounts that match. I review, accept, and the workbench opens with accounts already scored and ready for sellers.

Scoring Model — DC Methodology (Fit 50 · Need 35 · Intent 15)

Pillar	Weight	Key Dimensions
Fit	50%	Company size (bell-curved on offering ICP), industry tier match, IT spend bonus, complexity signals
Need	35%	Competitor installs (displacement), complementary tech presence, cross-sell signals, tech sophistication, stack momentum
Intent	15%	Direct product intent topics, category intent (tiered), competitor intent, TrustRadius/buyer activity

Default tier thresholds: $A \geq 75 \cdot B \geq 55 \cdot C \geq 35 \cdot D \geq 15$

Key Requirements

ID	Requirement
F3.1	Auto-generate one scoring model per confirmed offering on Step 3 entry (within 2 min).
F3.2	Every dimension is auto-populated from offering config (competitors → Displacement Targets; intent topics → Direct Intent; complementary
F3.3	Score all HG-universe accounts that pass the offering's basic ICP filter — not just the tenant's CRM book.
F3.4	Admin sees a tier distribution chart (A/B/C/D/Out counts) per offering.
F3.5	Admin can drill into top 20 A-tier accounts per offering for sanity check, with dimension-level score breakdown.
F3.6	Admin can adjust any dimension input and trigger a re-score.
F3.7	Admin must confirm each model before proceeding. Confirmed models become v1 with status 'active'.

Acceptance Criteria

- One scoring model per offering auto-generated within 2 min.
- All dimensions pre-populated from offering config.
- Initial scoring runs against HG-universe accounts matching offering ICP filter.
- Admin sees tier distribution + can view top 20 A-tier accounts with score breakdown.
- Admin can adjust dimensions and re-score.
- Admin cannot enter the workbench until all models are confirmed.

8. Step 4 — Workbench Entry + Enrichment

User Story: The moment I enter the workbench I see a filtered, scored list of accounts. I can ask it in plain language to add columns — 'show which are researching Wiz competitors', 'flag recent CISO hires' — and new columns appear with provenance.

Default Entry State

- **Offering lens:** filtered to first confirmed offering (lens switcher in header).
- **Account source:** HG-universe whitespace accounts (A + B tier) — the 'found accounts you didn't know about' moment.
- **Sort:** composite score descending.
- **Filters:** Tier A + B pre-applied (toggleable).

Default Columns (v1) — in display order

#	Column	Source	Notes
1	Account	name + domain	Click → Account view
2	Tier	accountScore.tier	Color chip: A green · B blue · C amber · D grey. Hover → breakdown.
3	Revenue	HG firmographics	Sortable
4	Employees	HG firmographics	Sortable
5	Why Now	Top scoring signals + intent for this offering	One-line summary (max 80 chars). Click → full provenance.
6	Competitive Insights	account.installs ∩ offering.competitors	Chips of competitor products installed. Hover → age, region.
7	Complementary Tech	account.installs ∩ offering.complementary	Chips of complementary stack. Indicates readiness.

Enrich with AI — Intent-Based Query Engine

A natural-language input in the toolbar accepts plain-language requests and returns new columns appended to the workbook, each cell with a provenance chip (source + refresh date).

Query example	Resulting column
'Show intent surge on Wiz CNAPP in last 30 days'	intent.cnapp_30d — surge level chip + topic list
'Add cloud provider for each account'	tech.cloud_provider — chip (AWS / Azure / GCP / Multi)
'Add cloud security spend last 12 months'	spend.cloud_security_12mo — \$ value + YoY trend
'Flag CISO or VP Security hires in last 90 days'	hiring.security_leader_90d — name + title + hire date
'Flag breach disclosures last 6 months'	events.breach_6mo — boolean + headline + date
'Show TrustRadius activity on Wiz competitors'	activity.competitor_trust_radius — date + activity type
'Add their renewal date' (if CRM connected)	crm.renewal_date — from Salesforce / HubSpot

Key Requirements

ID	Requirement
F4.1	Default columns render within 2s for up to 5,000 accounts on workbench entry.
F4.2	Offering lens switcher updates the table within ≤500ms (cached) / ≤3s (uncached).
F4.3	Source toggle (Book / Whitespace / Both) visible when CRM is connected.
F4.4	Enrich query engine accepts natural language, returns new column or clarifying follow-up.
F4.5	Every enriched cell has a provenance chip. Bulk enrichment for 5K accounts ≤ 30s target / 60s acceptable.
F4.6	Admin can pin enrichment columns to a Saved View. Enrichments are non-destructive.
F4.7	Queries requesting unlicensed data return 'Upgrade required' — not an error.

Acceptance Criteria

- Workbench renders default 7 columns in < 2s post-onboarding.
- Offering lens switcher and source toggle work as specified.
- 'Enrich with AI' input returns a new column with per-cell provenance.
- All 7 query categories return sensible columns.
- Admin can save enriched views; enrichment columns are removable with one click.
- Why Now column shows a meaningful one-liner for every A/B tier account.

9. Step 5 — Play Configuration

User Story: The platform proposes Plays based on my offerings — Competitive Takeout, Net New Logo, Expansion, Renewal Defense, High-Intent Active Buyer. Each is composed from 1P + 3P signals. I review, edit weights, and activate.

Concept: Plays vs. Signals

- **Play** = a business motion (Competitive Takeout, Net New Logo, Expansion, Renewal Defense, Active Buyer, Catalyst Event).
- **Signal** = an atomic, weighted, evaluable primitive that ranks accounts within a Play.
- A Play composes 3–7 signals with weights. Account rank = weighted sum of firing signals.
- Each Play has **recommended actions** (e.g., 'Send displacement email', 'Trigger competitive teardown workflow').

1P vs. 3P Signal Sources

Type	Signal examples	Dependency
1P (first-party)	CRM stage, last activity date, renewal date, opportunity status, owner	Requires CRM integration (Step 6)
3P (third-party)	Intent surge, competitor install/removal, complementary tech install, IT spend, third-party data	Requires third-party integrations (e.g., Clearbit, Bombora, TrustRadius)

Plays can be activated with 3P-only signals. 1P signals activate automatically when the relevant CRM integration connects in Step 6.

Key Requirements

ID	Requirement
F5.1	Propose 3–5 Plays on Step 5 entry, each tailored to confirmed offerings with offering-specific signal inputs.

ID	Requirement
F5.2	Each Play shows its 3–7 signals with default weights, labeled 1P or 3P. 1P signals without CRM show 'Awaiting integration'.
F5.3	Admin can add/remove signals, drag-to-reweight, and see a live ranked preview of the top 20 accounts.
F5.4	Each Play has recommended actions; admin can edit text or attach a workflow.
F5.5	Admin sets each Play's audience (AE / CSM / Both) and persona pin defaults.
F5.6	Admin must activate ≥ 1 Play to complete Step 5. Others can remain in 'draft'.
F5.7	Play evaluation reuses the runtime evaluator (playEvaluator.js) — no separate scoring pathway.

Acceptance Criteria

- 3–5 Plays appear pre-populated for the tenant's offerings.
- Each Play shows 1P + 3P signals with default weights.
- Live preview shows top 20 ranked accounts; updates within 1s on signal/weight change.
- Admin can activate, pause, or save-as-draft any Play.
- Plays requiring 1P signals show 'Awaiting integration' until Step 6 connects.
- Admin can complete Step 5 with ≥ 1 active Play.

10. Step 6 — Integrations Setup

User Story: I connect my CRM and Slack so the workbench can pull first-party data, sync enriched accounts back, and notify sellers when Plays match. I see exactly what data each integration accesses and which agents can write back.

Integrations (v1)

Integration	Type	Direction	Primary Purpose
Salesforce	CRM	Bi-directional	Book of accounts, opps, contacts, custom fields. Sync enriched accounts + scoring tiers back
HubSpot	CRM	Bi-directional	Same as Salesforce, for HubSpot-native tenants. Companies, deals, contacts, custom proper
Slack	Notification	Outbound + light inbound	Post Play matches, daily digest, alerts to channels or DMs.

CRM Setup Flow (Salesforce + HubSpot)

- OAuth connect (Salesforce: pick sandbox vs production).
- Object scope selection: Accounts (required), Contacts, Leads, Opportunities, Tasks/Activities, Notes, Users/Roles.
- Field mapping: auto-map standard fields; admin reviews + approves custom field mappings.
- Sync direction per object (defaults: Accounts bi-directional; Tasks/Notes write-only; Users read-only).
- Agent write-back scopes — **all off by default**; admin explicitly enables each:
 - `crm.account.create / update` — create whitespace as CRM accounts; enrich existing accounts
 - `crm.contact.create / update` — export discovered contacts; enrich existing contacts
 - `crm.opportunity.update` — update opp fields from agent suggestions
 - `crm.task.create / crm.note.create` — log agent actions + briefs
 - `crm.bulk.import.scores` — bulk-write scoring tiers + Why-Now to Account custom fields
- Test sync (single-account read + test write + pagination check) must pass before integration is marked active.
- Initial bulk pull (Accounts + Contacts + Opportunities + Users) → admin notified when complete.

Integration-Complete Notification

When bulk pull completes, the system notifies the admin on all channels in parallel: **in-app banner** (with CTA 'Add Sellers →'), **email** (synced object counts + enabled scopes), and **Slack DM** (if Slack is connected). The notification enumerates every enabled and available-but-disabled agent scope in a single permissions ledger. Step 7 (Configure Teams) unlocks immediately after the first CRM integration completes.

Key Requirements

ID	Requirement
F6.1	OAuth flows for Salesforce, HubSpot, and Slack all work from Step 6.
F6.2	Onboarding does NOT require any integration to complete (graceful whitespace-only path).
F6.3	All agent write-back scopes default to off; require explicit admin toggle.
F6.4	Test sync must pass before integration status becomes 'active'.
F6.5	Integration-complete notification fires within 5 min, on all configured channels.

ID	Requirement
F6.6	Notification enumerates exact synced object counts + every enabled/disabled agent scope.
F6.7	Step 7 unlocks immediately after first CRM integration completes.
F6.8	Connecting CRM mid-onboarding re-evaluates 1P signals on existing Plays within 30s.
F6.9	Disconnect is one-click, reversible; data retained.

Acceptance Criteria

- Admin can connect Salesforce, HubSpot, and Slack via OAuth from Step 6.
- Admin can complete onboarding without connecting any integration.
- All agent write-back scopes default to off and require explicit toggle.
- Test sync passes before integration becomes 'active'.
- Integration-complete notification fires within 5 min on all configured channels.
- Connecting CRM re-evaluates 1P signals on existing Plays within 30s.

11. Step 7 — Configure Teams

User Story: I organize sellers into Teams, each with its own offerings, scoring profile, and default Plays. A seller's workbench is pre-configured from their team. Without teams, every seller inherits tenant-wide defaults — too coarse to be useful.

Teams are the inheritance layer between tenant-level config and per-seller config: **Tenant** → **Team** → **Seller**. When a seller logs in, their workbench is pre-filtered to their team's offerings, pre-scored with their team's scoring profile, and pre-loaded with their team's pinned Plays.

Per-Team Configuration

Field	Source / Options	Required?
Name	Admin-entered (e.g., 'Enterprise West', 'Mid-Market AMER')	Yes
Offerings[]	Select from confirmed offerings (Step 2); ≥1 required	Yes
Primary Scoring Profile	Pick from scoring models (Step 3); exactly 1	Yes
Default Plays[]	Select from active Plays (Step 5)	Yes (≥1)
Audience	AE · AM · CSM · BDR/SDR · Mixed	Yes
Territory	Optional CRM territory mapping (drives auto-seller suggestion in Step 9)	No
Manager	Optional RGI user id	No

Key Requirements

ID	Requirement
F7.1	Step 7 unlocks after Step 5 (Plays) is confirmed.
F7.2	System proposes a default 'All Sellers' team on entry, plus one team per offering if ICPs are distinct.
F7.3	No cap on number of teams during onboarding.

ID	Requirement
F7.4	Admin can clone a team (for territory variants like 'Enterprise East').
F7.5	Admin can delete a team only if no sellers are assigned; otherwise must force-reassign.
F7.6	Changes to a team's offerings/Plays propagate to all assigned sellers on next refresh.
F7.7	Live preview shows account count + A/B tier breakdown per team's scoring profile.
F7.8	Admin must activate ≥1 Team before proceeding to Step 8.

Acceptance Criteria

- Admin sees 1+ proposed teams on Step 7 entry.
- Admin can create, edit, clone, delete, and activate teams.
- Each team has ≥1 offering, exactly 1 primary scoring profile, ≥1 default Play.
- Preview shows realistic book size + A/B tier counts per team.
- Changes to team config propagate to assigned sellers without manual refresh.

12. Step 8 — Configure Agents

User Story: Before I invite sellers, I decide which AI agents they can use. I enable them at the Team level so sellers in different motions get the right toolset. Enabled agents appear in the seller's Account view as one-click affordances.

Agent Catalog (v1)

Agent	Purpose	Required Scope	Default Visibility	Cost/Call
Account Brief	Multi-section account briefing: firmographics, tech stack, brief generates (tailored, contextual, and data-driven) next steps	Account	On	\$0.05
Email Outreach	Draft personalized outreach: cold, follow-up, disengagement, and prospect follow-up	Account, Sales Enablement	On	\$0.02
Opportunity Finder	Scans account(s) for opportunity signals: intent signals, competitive scan (oversight), lead discovery, funding	Account, Sales Enablement	On	\$0.10
Find More Contacts	Discovers buying-committee contacts not in CRM	Account, Sales Enablement	On	\$0.08

Configuration Granularity

Level	What's configured	When to use
Tenant	Which agents are available at all (kill switch)	Set once; rarely changed
Team	Which available agents are enabled per team (inherits audience-aware defaults)	Primary use — configure during Step 8
Seller (override)	Enable/disable specific agents vs. team default	Pilot users, training cohorts, special accounts

Key Requirements

ID	Requirement
F8.1	Step 8 unlocks after Step 7 (Teams) is complete.
F8.2	Admin sees all 4 agents with description, cost profile, required + optional scopes, and default visibility.
F8.3	For each agent: admin can toggle tenant-wide availability (kill switch).
F8.4	Team agent matrix: rows = teams, columns = agents. Each cell is a toggle. Audience-aware defaults applied on first entry.
F8.5	Admin can grant optional scopes per team via a scope drawer.
F8.6	Cost summary sidebar shows estimated cost per seller per day based on enabled agents.
F8.7	'See seller view' preview shows exactly what the seller's Account view will look like.
F8.8	Disabling an agent removes its affordance from sellers' views within 60s.
F8.9	Admin can complete Step 8 with zero agents enabled (graceful path).

Acceptance Criteria

- Admin sees all 4 agents with descriptions, scopes, and costs.
- Tenant-wide availability and per-team enablement toggles work.
- Audience-aware defaults apply on first entry.
- Cost summary updates live as settings change.
- Preview shows exact seller Account view.
- Disabling an agent removes it from sellers' views within 60s.

13. Step 9 — Invite Sellers

User Story: With Teams configured and my CRM connected, the system reads my CRM users, suggests which team each seller belongs to, and lets me bulk-assign and invite. Sellers click a magic link and land in their pre-configured workbench.

Step 9 unlocks after: CRM integration (Step 6) completes + ≥ 1 Team active (Step 7) + agent configuration saved (Step 8).

Seller Discovery + Auto-Assignment Logic

- Pull all active CRM Users / Owners with role matching Sales, AE, AM, CSM, BDR, SDR.
- **Territory match:** if CRM territory matches a Team's territory hints → suggest that team.
- **Role match:** if no territory match, suggest team whose audience aligns with seller's role.
- **Fallback:** assign to default 'All Sellers' team.

Per-Seller Config (all editable by admin)

Field	Default	Admin can override?
Team	Auto-assigned by territory/role logic	Yes (drag-drop or picker)
Offering lens	Team's primary offering	Yes (per-seller override)
Pinned Plays	Team's default Plays	Yes (add/remove per-seller)
Enabled agents	Team's agent config	Yes (enable/disable per-seller)
Persona	Inferred from CRM role	Yes (AE / AM / CSM / BDR / Admin)
Slack handle	Matched from Slack workspace	Yes

Key Requirements

ID	Requirement
F9.1	System surfaces CRM-detected seller list within 30s of Step 9 entry.
F9.2	Admin can bulk-select, filter by role/territory/team, and assign/reassign in batch.
F9.3	Invitation triggers: templated email (magic link + team + book size preview + enabled agents) + Slack DM if connected.
F9.4	Magic link is single-use, expires in 14 days, tied to the inviting tenant.
F9.5	Seller status lifecycle: proposed → invited → pending_acceptance → active.
F9.6	Admin can impersonate-view any seller to see exactly what they'll see post-acceptance.
F9.7	Bulk invite of up to 500 sellers completes within 60s.
F9.8	Onboarding marks 'complete' once ≥ 1 seller is active.

Acceptance Criteria

- Step 9 unlocks only after CRM (Step 6), ≥ 1 Team (Step 7), and agent config (Step 8).
- Seller list pulled from CRM within 30s.
- Each seller pre-assigned to a team; admin can override.

-
- Bulk invite of up to 500 sellers works; invites delivered via email + Slack DM.
 - Impersonate-view shows exact seller experience.
 - Onboarding marks complete once ≥ 1 seller is active.

14. Step 10 — Seller First-Run Experience

User Story: I click the magic link, set up my account (< 2 min), and land directly in the workbench. My book is already there — scored, with Why-Now reasons, competitor insights, and complementary tech. I also see whitespace I didn't know about. A 30-second tour, then I'm working.

Entry Path

Step	What seller does	Time
Welcome	Team name + 'Get started' CTA. No marketing fluff.	5s
Auth	SSO (if tenant has SAML/OAuth) or password + optional MFA.	30–60s
Profile	Display name + time zone (auto-detected). Slack handle pre-filled.	15s
Workbench	Auto-routed. No intermediate dashboard.	—

Pre-Loaded Workbench

What's loaded	Details
Book of accounts	CRM-owned accounts, scored against team's primary scoring profile. Sort: composite score desc. Filter: A+B tier (top)
Whitespace accounts	HG-universe accounts matching team's ICP, not in seller's CRM, scoring ≥ Tier C. Source toggle: Book / Whitespace
Default columns	Account · Tier · Revenue · Employees · Why Now · Competitive Insights · Complementary Tech (scoped to seller's c
Plays sidebar	Team's default Plays pinned with match-count badges. Click → filters workbook to ranked accounts.
Agents (Account view)	Enabled agents appear as one-click buttons in the Account header. Brief · Email · Opportunities · Contacts.

Orientation Tour (First Entry Only, Skippable)

#	UI Element	Copy hint
1	Workbook table	'Your book — scored against [Offering]. Why Now tells you who matters today.'
2	Source toggle	'You also have whitespace — companies HG knows about not in your CRM.'
3	Plays sidebar	'Your team focuses on these Plays. Click one to see ranked accounts.'
4	Account row → Account view	Click any account. [N] agents are ready: brief, email, contacts, opportunities.'
5	Enrich + Activate	'Add columns in plain language, or run an agent. You're set.'

Key Requirements

ID	Requirement
F10.1	Magic link validates token, creates RGI user account, routes to account setup.
F10.2	Expired/consumed tokens show clear error with admin contact info.
F10.3	Account setup supports SSO (OIDC + SAML) and password + optional MFA.

ID	Requirement
F10.4	Workbench pre-loads book + whitespace + scores + Plays; first-paint ≤ 2s.
F10.5	Orientation overlay shows on first entry only (dismissed state persists per seller). Re-triggerable from '?' icon.
F10.6	First-day drawer surfaces top 3 A-tier book accounts + top 3 whitespace + most-active Play (all with agent affordances scoped to seller's pe
F10.7	Telemetry captures: time invite-sent → first login, time first login → first action, orientation completion rate.

Acceptance Criteria

- Magic link → setup → workbench flow works end-to-end.
- Expired/invalid tokens show a clear error.
- Account setup supports both SSO and password.
- Workbench loads with book + whitespace + scores + Plays in ≤ 2s.
- Orientation overlay shows on first entry only; re-triggerable.
- First-day drawer surfaces 3 A-tier + 3 whitespace + most-active Play.

15. Ongoing Seller Workspace

Design principle: Default to act, not browse. Every account row has quick-action affordances. Agent runs land in the same view. Every enriched cell and agent output has one-hover provenance. Persistence over rediscovery.

Workspace Anatomy

Surface	What it does
Plays Sidebar (left)	Pinned Plays with match-count badges. Saved Views below. Click Play → filters workbook.
Workbook Table (center)	Book / Whitespace / Both. 7 default columns + enriched columns. Multi-sort, filter chips.
Account View (on row-click)	Agent action row (Brief · Email · Opportunities · Contacts) + tabs: Overview, Signals, Contacts, Activity.
Header	Offering lens switcher, global search, help / orientation re-trigger.

Whitespace → Book Conversion

- Seller finds a whitespace account → 'Add to book' button.
- If `crm.account.create` scope is granted: account pushed to CRM with seller as Owner. Flips from whitespace to book in real time.
- If scope not granted: account added to an RGI-local Watchlist; admin can batch-approve CRM push.

Account View — Agent Action Row

The agent action row is the first thing the seller sees after the account header. Clicking opens a small popover (1–3 fields), fires the agent run, and streams output into the activity feed — no page navigation. Outputs are timestamped with agent provenance + cost and are editable/regeneratable.

Agent	Popover inputs	Output
Account Brief	Brief type (Discovery / QBR / Competitive) + offering context	Streaming Brief activity feed. One-click 'Export to CRM as Note'.

Agent	Popover inputs	Output
Email Outreach	Email type + contact target + optional tone hint	Draft in activity feed. 'Send' (if scope enabled) or 'Copy'. 'Log as Task'.
Opportunity Finder	Time window (default 90d) + hypothesis type filter	Ranked hypothesis cards with evidence + recommended action chain.
Find More Contacts	Target persona (from offering's buying committee) + account	Contact list. 'Add to CRM' per contact if scope enabled.

16. Cross-Cutting Requirements

Persistence & Resumability

- Every step's state persists immediately. Admin can close browser and resume.
- A resume prompt appears on every login until all admin steps (1–8) are confirmed or skipped.
- Each step has a 'Save & exit' affordance.

AI Agent Architecture

- Each step driven by a dedicated agent type (Tenant Context, Offering Grouping, Scoring Model, Enrichment, Play Proposal, Team Suggestion, Seller Role Mapping).
- All agents are **idempotent** (same inputs → same outputs) and **observable** (every run logs prompt, model version, raw output, final decision).
- Agents chain: Step 2 depends on Step 1; Step 5 depends on Steps 2 + 3; Step 7 depends on offerings + Plays; Step 9 depends on Step 6 CRM users + Step 7 Teams.

Error States

Error type	Handling
Agent failure	Friendly retry. If repeated, offer manual entry fallback.
HG data gap	Gracefully degrade to manual-entry for affected fields.
Integration auth failure	Clear error + reconnect CTA; admin can skip + revisit.
Magic link expired/consumed	Clear seller-facing error with admin contact info.
Invitation bounce	Per-seller 'Re-send' / 'Edit email' affordance.

Performance Budgets

Step	Target	Acceptable
Step 1 — Tenant extraction	≤ 60s	—
Step 2 — Offering generation	≤ 30s	—
Step 3 — Model generation per offering	≤ 2 min	—
Step 3 — Initial universe scoring	≤ 10 min	Background after model confirmation
Step 4 — Workbench first paint (≤5K accounts)	≤ 2s	—
Step 4 — Enrichment query (≤5K accounts)	≤ 30s	≤ 60s
Step 5 — Play proposal	≤ 30s	—
Step 5 — Live preview update on signal change	≤ 1s	—
Step 6 — OAuth + field mapping load	≤ 10s	—
Step 9 — Bulk invite 500 sellers	≤ 60s	—
Step 10 — Seller workbench first paint	≤ 2s	—

Security & Privacy

- Tenant context is tenant-private by default. No cross-tenant raw data exposure.
- Integration credentials encrypted at rest; OAuth tokens never logged in plaintext.
- Invitation magic links are single-use, time-bound (14 days), tenant-scoped.
- Sellers can only see accounts within their team's offering scope; cross-team visibility is admin-gated.
- All agent scope grants logged in the agent-access audit log (existing infrastructure).

17. Success Metrics

Metric	Target	Why it matters
Time to confirmed workbench (Steps 1–3)	< 60 min (median)	Core TTV claim
Time to fully-configured tenant (Steps 1–8)	< 150 min (median)	Admin onboarding cost
Step 1 field edit rate	< 30% of fields edited	Extraction precision
Step 2 offering edit rate	< 50% of attributes edited	LLM grouping quality
Step 3 model approval (no dimension edits)	> 80%	Model auto-build quality
Step 4 enrichment query success rate	> 75% return usable column	Query engine quality
Step 5 Play activation rate	≥1 active Play for > 80% of tenants	Play proposal relevance
Step 6 CRM connection rate (within 30 days)	> 70%	1P data flow health
Step 9 seller acceptance (within 7 days)	> 60%	Invitation UX quality
Step 10 first-action latency	< 10 min from first login	Workbench immediate usability
Step 10 first agent invocation rate	> 50% within first session	Agent discoverability + trust
Admin resume rate (started + didn't finish)	< 20%	Overall onboarding friction

18. Phasing & Rollout

Phase	Scope	Target
A	Step 1 — Tenant Context Extraction (end-to-end + manual fallback)	Q3
B	Step 2 — Offering Configuration (LLM grouping)	Q3 / early Q4
C	Step 3 — Scoring Models + Initial Scoring (DC scoring SQL generators)	Q4
D	Step 4 — Workbench Entry + intent-based enrichment query engine	Q4
E	Step 5 — Play Configuration (proposal + signal composition + live preview)	Q4 / early Q5
F	Step 6 — Integrations (Salesforce + HubSpot + Slack + completion notification)	Q5
G	Step 7 — Teams (data model + configurator + inheritance)	Q5
H	Step 8 — Agents (4 v1 agents + visibility + scope grants + cost telemetry)	Q5 / early Q6
I	Step 9 — Seller Invitation (CRM discovery + bulk invite + magic link)	Q6
J	Step 10 — Seller First-Run (setup + workbench pre-load + orientation + agents)	Q6
K	Stitching + polish (full 10-step, resume, telemetry, error UX, E2E QA)	Q6 / early Q7

19. Open Questions

#	Question	Recommendation
1	What's the confidence threshold for blocking confirmation (Step 1)?	> 0.6 = visual flag; < 0.3 = block confirmation until reviewed
2	Should we ask for admin email to validate domain ownership?	Trust domain unverified in v1
3	Tenant has only 1–2 products — pad to 3 offerings or accept fewer?	Accept fewer; don't force artificial offerings
4	Should Step 3 allow admins to skip entirely (auto-accept model, skip in the background)?	More in the background; onboarding moment is too valuable to skip
5	Primary CRM if tenant has both Salesforce + HubSpot?	Admin picks primary; secondary is read-only for enrichment context
6	Should sellers be able to request access to a disabled agent in-app?	Out-of-band (ask admin) in v1; in-app request flow in v2
7	Rate limits on agents — ship unlimited in v1 or per-seller daily caps?	Unlimited in v1; add caps post-MVP after usage data informs the right threshold
8	Orientation overlay — mandatory or always-dismissible?	Always-dismissible (sellers hate forced flows)
9	Magic link expiry — 14 days or shorter?	14 days is reasonable; security review may shorten
10	For the 'Both' workbook source — visually distinguish book vs. whitespace rows?	Yes, pad the edge on whitespace rows so sellers always know which world a row

20. Glossary

Term	Definition
Tenant	The HG customer company (e.g., Wiz.io).
Offering	A distinct product/solution the tenant sells (e.g., Wiz CNAPP). Max 3 during onboarding.
Scoring Model	DC methodology model scoring accounts on Fit / Need / Intent for a given offering.
Play	A business motion (Competitive Takeout, Net New, Expansion, Renewal Defense, Active Buyer) composed from weighted
Signal (Ranking)	Atomic, weighted, evaluable primitive that ranks accounts within a Play. 1P (from CRM) or 3P (from HG).
HG-universe account	Any company HG has data on. Distinct from a 'book account' in the tenant's CRM.
Whitespace	HG-universe accounts that score well for an offering but are NOT in the tenant's CRM book.
DC methodology	Doug Cardamone's scoring framework: Fit 50% · Need 35% · Intent 15%.
RGI Workbench	The seller's primary workspace. Book + whitespace + Plays + agents.
Enrichment column	A workbook column added via the intent-based query engine, with per-cell provenance.
Team	Grouping of sellers sharing offerings, scoring profile, default Plays, and enabled agents. Inheritance layer between tenant
Agent	Seller-facing AI capability invocable in the Account view (Brief, Email, Opportunity Finder, Find More Contacts). Visibility c
Provenance	Visible source + reason for every AI-extracted or AI-derived value. Builds trust.
1P signal	First-party signal from the tenant's connected CRM (stage, activity date, renewal date, etc.).
3P signal	Third-party signal from HG data (intent surge, competitor install, spend trend, hiring, news).